Direct cross-family receiver prediction has a small, heterogeneous mean gain

Held density gain vs zero (positive better)

Native training
Reversed training polarity
Circular75s training shift
Untrained prior covariance

Projected differential RMS change (Hz)

0010 0011 0012 0013 0014 0015 0016 0017 0018 0019 0075 0077 0079 0081 0083 0085 0087 0089 0091 0093 0095 0097 0099 0101

Same held affine-null vector, complete Gaussian covariance and normalization for every predictor.
Whole-family leave-outs overlap in training; summed gains are diagnostics, not independent significance or hardware evidence.